

Training loss of a 2-layer network (backprop gradients)

